


# Chapter 4


## Interval scheduling



Slides by Kevin Wayne.  
Copyright © 2005 Pearson-Addison Wesley.  
All rights reserved.

# Interval scheduling

- Job  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Two jobs are **compatible** if they don't overlap.
- Goal: find **maximum** subset of mutually **compatible** jobs.





# Interval scheduling

---

- **Input:**
  - A set of  $n$  intervals  $I_1, \dots, I_n$
  - interval  $I_i$  has starting time  $s_i$  and finish time  $f_i$
- **Feasible solution:**
  - A subset  $S$  of the intervals that are mutually compatible, i.e. for each  $I_i, I_j \in S$ ,  $I_i$  does not overlap with  $I_j$
- **Measure (to maximize):**
  - number of scheduled intervals, i.e. cardinality of  $S$





# Interval scheduling: greedy algorithms

---

**Greedy template.** Consider jobs in some natural order.

Take each job provided it's compatible with the ones already taken.

-  **[Earliest start time]** Consider jobs in ascending order of  $s_j$ .  

- **[Earliest finish time]** Consider jobs in ascending order of  $f_j$ .  

-  **[Shortest interval]** Consider jobs in ascending order of  $f_j - s_j$ .
-  **[Fewest conflicts]** For each job  $j$ , count the number of conflicting jobs  $c_j$ . Schedule in ascending order of  $c_j$ .

DEFINIZIONE DI ALGORITMO GREEDY:

UNA SERIE DI DECISIONI LOCALI OTTIME CHE PORTANO (FORSE) A UNA SOLUZIONE GLOBALE OTTIMA

# Interval scheduling: greedy algorithms

---

**Greedy template.** Consider jobs in some natural order.  
Take each job provided it's compatible with the ones already taken.

**counterexample for earliest start time**



**counterexample for shortest interval**



**counterexample for fewest conflicts**





# Interval scheduling: earliest-finish-time-first algorithm

---

EARLIEST-FINISH-TIME-FIRST ( $n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$ )

---

**Sort** jobs by finish times and renumber so that  $f_1 \leq f_2 \leq \dots \leq f_n$ .

$S \leftarrow \emptyset$ .  $\longleftarrow$  set of jobs selected

**FOR**  $j = 1$  **TO**  $n$

**IF** (job  $j$  is compatible with  $S$ )



$S \leftarrow S \cup \{ j \}$ .

**RETURN**  $S$ .

---

# Earliest-finish-time-first algorithm demo


---






# Earliest-finish-time-first algorithm demo

---





**job B is compatible (add to schedule)**






# Earliest-finish-time-first algorithm demo


---



# Earliest-finish-time-first algorithm demo





**job C is incompatible (do not add to schedule)**






# Earliest-finish-time-first algorithm demo


---



# Earliest-finish-time-first algorithm demo





**job A is incompatible (do not add to schedule)**






# Earliest-finish-time-first algorithm demo


---



# Earliest-finish-time-first algorithm demo





**job E is compatible (add to schedule)**






# Earliest-finish-time-first algorithm demo



# Earliest-finish-time-first algorithm demo





**job D is incompatible (do not add to schedule)**






# Earliest-finish-time-first algorithm demo

---






# Earliest-finish-time-first algorithm demo





**job F is incompatible (do not add to schedule)**




# Earliest-finish-time-first algorithm demo

---







# Earliest-finish-time-first algorithm demo




**job G is incompatible (do not add to schedule)**



# Earliest-finish-time-first algorithm demo





**job G is incompatible (do not add to schedule)**





# Earliest-finish-time-first algorithm demo


---






# Earliest-finish-time-first algorithm demo

---





**job H is compatible (add to schedule)**



# Earliest-finish-time-first algorithm demo


---






# Earliest-finish-time-first algorithm demo

---



done (an optimal set of jobs)



# Interval scheduling: earliest-finish-time-first algorithm

**EARLIEST-FINISH-TIME-FIRST** ( $n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$ )

**SORT** jobs by finish times and renumber so that  $f_1 \leq f_2 \leq \dots \leq f_n$ .

$S \leftarrow \emptyset$ .  $\leftarrow$  set of jobs selected

**FOR**  $j = 1$  **TO**  $n$

**IF** (job  $j$  is compatible with  $S$ )

$S \leftarrow S \cup \{ j \}$ .


**RETURN**  $S$ .

**Proposition.** Can implement earliest-finish-time first in  $O(n \log n)$  time.

- Keep track of job  $j^*$  that was added last to  $S$ .

- Job  $j$  is compatible with  $S$  iff  $s_j \geq f_{j^*}$ .  $\leftarrow$  INIZIANDO VOI  $\geq$  FINE VECCHIO

- Sorting by finish times takes  $O(n \log n)$  time.





# Interval scheduling: analysis of earliest-finish-time-first algorithm

---

Let  $i_1, i_2, \dots, i_k$  be set of jobs selected by greedy (ordered by finish times).


Let  $j_1, j_2, \dots, j_m$  be set of jobs in an optimal solution (ordered by finish times)

denote by  $f(i_r)$  the finish time of job  $i_r$ .

**Lemma.** For every  $r=1,2,\dots,k$ , we have  $f(i_r) \leq f(j_r)$ .

**Pf.** [by induction]

- $r=1$ : Obvious.
- $r>1$ :






# Interval scheduling: analysis of earliest-finish-time-first algorithm


**Theorem.** The earliest-finish-time-first algorithm is optimal.

**Pf.** [by contradiction]

- Let  $i_1, i_2, \dots, i_k$  be set of jobs selected by greedy (ordered by finish times).
- Let  $j_1, j_2, \dots, j_m$  be set of jobs in an optimal solution (ordered by finish times)
- Assume greedy is not optimal
- then  $m > k$

ABBIAMO UNA CONTRADDIZIONE  
IL GREEDY LO AVREBBE SELEZIONATO





## SECTION 4.1

# 4. GREEDY ALGORITHMS I

---

*A related problem:*


- *interval partitioning*



# Interval partitioning

- Lecture  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

Ex. This schedule uses 4 classrooms to schedule 10 lectures.






# Interval partitioning

- Lecture  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

Ex. This schedule uses 3 classrooms to schedule 10 lectures.



# Interval partitioning

---

- **Input:**

- A set of  $n$  intervals  $I_1, \dots, I_n$
- interval  $I_i$  has starting time  $s_i$  and finish time  $f_i$

- **Feasible solution:**

- A partition of the intervals into subsets (called classrooms)  $C_1, \dots, C_d$  such that each  $C_i$  contains mutually compatible intervals

- **Measure (to minimize):**

- number of classrooms, i.e.  $d$



## Interval partitioning: greedy algorithms

---

**Greedy template.** Consider lectures in some natural order. Assign each lecture to an available classroom (which one?); allocate a new classroom if none are available.

- [Earliest start time] Consider lectures in ascending order of  $s_j$ .
- [Earliest finish time] Consider lectures in ascending order of  $f_j$ .
- [Shortest interval] Consider lectures in ascending order of  $f_j - s_j$ .
- [Fewest conflicts] For each lecture  $j$ , count the number of conflicting lectures  $c_j$ . Schedule in ascending order of  $c_j$ .




# Interval partitioning: greedy algorithms

---

**Greedy template.** Consider lectures in some natural order. Assign each lecture to an available classroom (which one?); allocate a new classroom if none are available.


**counterexample for earliest finish time**



**counterexample for shortest interval**



**counterexample for fewest conflicts**





## Interval partitioning: earliest-start-time-first algorithm

---

**EARLIEST-START-TIME-FIRST** ( $n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$ )

**SORT** lectures by start times and renumber so that  $s_1 \leq s_2 \leq \dots \leq s_n$ .

$d \leftarrow 0$ .  number of allocated classrooms

**FOR**  $j = 1$  **TO**  $n$

**IF** (lecture  $j$  is compatible with some classroom)

        Schedule lecture  $j$  in any such classroom  $k$ .

**ELSE**

        Allocate a new classroom  $d + 1$ .

        Schedule lecture  $j$  in classroom  $d + 1$ .

$d \leftarrow d + 1$ .

**RETURN** schedule.


---

# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.






# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**no compatible classroom: open up a new classroom and assign lecture to it**






# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**no compatible classroom: open up a new classroom and assign lecture to it**




# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**no compatible classroom: open up a new classroom and assign lecture to it**






# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture d is compatible with classrooms 1 and 3**




# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture e is compatible with classroom 1**






# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture f is compatible with classroom 2 and 3**






# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture g is compatible with classroom 2**




# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture h is compatible with classroom 1**






# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture j is compatible with classrooms 2 and 3**




# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture i is compatible with classroom 2**






# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

done





## Interval partitioning: earliest-start-time-first algorithm

---

**Proposition.** The earliest-start-time-first algorithm can be implemented in  $O(n \log n)$  time.

**Pf.**

- Sorting by start times takes  $O(n \log n)$  time.
- Store classrooms in a **priority queue** (key = finish time of its last lecture).
  - to allocate a new classroom, INSERT classroom onto priority queue.
  - to schedule lecture  $j$  in classroom  $k$ , INCREASE-KEY of classroom  $k$  to  $f_j$ .
  - to determine whether lecture  $j$  is compatible with any classroom, compare  $s_j$  to FIND-MIN
- Total # of priority queue operations is  $O(n)$ ; each takes  $O(\log n)$  time. ▪

CODA CON PRIORITA'  
TIPO DIZIONARIO

**Remark.** This implementation chooses a classroom  $k$  whose finish time of its last lecture is the **earliest**.

## Interval partitioning: lower bound on optimal solution


---

**Def.** The **depth** of a set of open intervals is the maximum number of intervals that contain any given point.

**Key observation.** Number of classrooms needed  $\geq$  depth.

**Q.** Does minimum number of classrooms needed always equal depth?

**A.** Yes! Moreover, earliest-start-time-first algorithm finds a schedule whose number of classrooms equals the depth.





## Interval partitioning: analysis of earliest-start-time-first algorithm

---

**Observation.** The earliest-start-time first algorithm never schedules two incompatible lectures in the same classroom.

**Theorem.** Earliest-start-time-first algorithm is optimal.

**Pf.**

- Let  $d$  = number of classrooms that the algorithm allocates.
- Classroom  $d$  is opened because we needed to schedule a lecture, say  $j$ , that is incompatible with a lecture in each of  $d - 1$  other classrooms.
- Thus, these  $d$  lectures each end after  $s_j$ .
- Since we sorted by start time, each of these incompatible lectures start no later than  $s_j$ .
- Thus, we have  $d$  lectures overlapping at time  $s_j + \epsilon$ .
- Key observation  $\Rightarrow$  all schedules use  $\geq d$  classrooms. ▪



## Dimostrazione di ottimalità

Vogliamo dimostrare che **questo algoritmo usa il numero minimo possibile di aule.**

### 1. Definiamo $(d)$ come il numero di aule utilizzate dall'algoritmo.

- Questo significa che esiste almeno un momento nel tempo in cui ci sono **esattamente  $(d)$  lezioni sovrapposte.**

### 2. Perché è stata aperta un'aula $(d)$ ?

- Significa che la lezione  $(j)$  assegnata a questa aula è **incompatibile con le lezioni nelle altre  $(d-1)$  aule.**
- In altre parole, ci sono  $(d)$  lezioni che si **sovrappongono** in un certo momento.

### 3. Tutte queste $(d)$ lezioni finiscono dopo l'inizio di $(j)$ .

- Poiché l'algoritmo ordina le lezioni per **orario di inizio crescente**, sappiamo che queste  $(d)$  lezioni iniziano **al più tardi** all'inizio di  $(j)$ , cioè non possono iniziare dopo.

### 4. A questo punto, abbiamo $(d)$ lezioni che si sovrappongono in un certo istante $(s_j + \epsilon)$ (subito dopo l'inizio di $(j)$ ).

- Questo significa che **qualunque altro algoritmo che cerchi di minimizzare il numero di aule avrà comunque bisogno di almeno  $(d)$  aule** per gestire queste lezioni senza conflitti.

### 5. Conclusione: l'algoritmo greedy usa esattamente $(d)$ aule, che è il minimo possibile.

- **Qualsiasi altro algoritmo dovrà usare almeno  $(d)$  aule**, perché non è possibile avere meno di  $(d)$  aule con  $(d)$  lezioni che si sovrappongono.
- Quindi, l'algoritmo greedy è **ottimale**.